

Model Evaluation

Classification quality metrics, prediction error, cross-validation

Machine Learning and Data Mining, 2024

Majid Sohrabi

National Research University Higher School of Economics



October 02, 2024

Classification quality metrics




How to evaluate a classifier?

0-1 Loss

Probability of an error (error rate):

$$\mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(y_i \neq \hat{y}_i)$$



0-1 Loss

Probability of an error (error rate):


$$\mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(y_i \neq \hat{y}_i)$$

Accuracy:

$$accuracy = 1 - \mathcal{L}_{0-1} = \frac{1}{N} \sum_{i=1 \dots N} \mathbb{I}(y_i = \hat{y}_i)$$

Not always a good quality measure

- E.g. when classes are imbalanced



Confusion matrix

		Actual class			
		1	2	...	C
Predicted class	1	n_{11}	n_{12}	\cdots	n_{1C}
	2	n_{21}	n_{22}	\cdots	n_{2C}
	\vdots	\vdots	\vdots	\ddots	\vdots
	C	n_{C1}	n_{C2}	\cdots	n_{CC}

n_{ij} – number of objects of class j , that were predicted as class i

Diagonal elements – correct classifications

Off-diagonal elements – incorrect classifications

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{True positive rate, TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

$$\text{True positive rate, TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{False positive rate, FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

$$\text{True positive rate, TPR} = \frac{TP}{TP+FN}$$

$$\text{False positive rate, FPR} = \frac{FP}{FP+TN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

$$\text{True positive rate, TPR} = \frac{TP}{TP+FN}$$

$$\text{False positive rate, FPR} = \frac{FP}{FP+TN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

$$\text{Recall} = \frac{TP}{TP+FN}$$

Binary case

		Actual class	
		+	-
Predicted class	+	TP (True Positives)	FP (False Positives)
	-	FN (False Negatives)	TN (True Negatives)

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

$$\text{True positive rate, TPR} = \frac{TP}{TP+FN}$$

$$\text{False positive rate, FPR} = \frac{FP}{FP+TN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

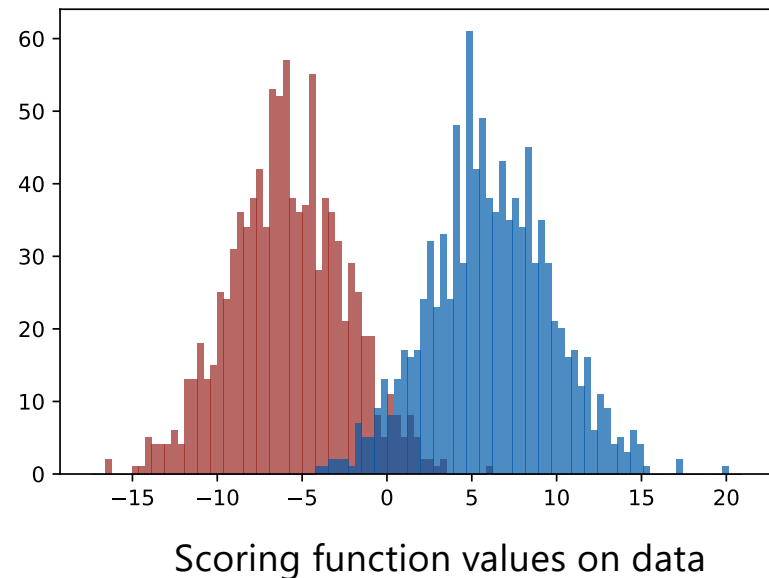

$$\text{Recall} = \frac{TP}{TP+FN}$$

$$F_1\text{-score} = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Continuous predictions

Many classification algorithms work with continuous scoring functions



- E.g. log odds in Logistic Regression, or scoring function of an SVM model



Continuous predictions


Many classification algorithms work with continuous scoring functions

- E.g. log odds in Logistic Regression, or scoring function of an SVM model



ROC-curve

Receiver operating characteristic = TPR as a function of FPR



History [\[edit \]](#)

The ROC curve was first used during [World War II](#) for the analysis of [radar signals](#) before it was employed in [signal detection theory](#).^[45] Following the [attack on Pearl Harbor](#) in 1941, the United States army began new research to increase the prediction of correctly detected Japanese aircraft from their radar signals. For these purposes they measured the ability of a radar receiver operator to make these important distinctions, which was called the Receiver Operating Characteristic.^[46]

https://en.wikipedia.org/wiki/Receiver_operating_characteristic

Nice demo: <http://arogozhnikov.github.io/2015/10/05/roc-curve.html>

ROC AUC probabilistic interpretation

ROC AUC = area under the ROC curve

For the population distribution:


$$P(x, y), \quad x \in \mathbb{R}^d, \quad y \in \{0, 1\}$$
$$\hat{f}(x): \mathbb{R}^d \rightarrow \mathbb{R} \quad - \text{classifier scoring function}$$

ROC AUC also equals the probability that

$$P[\hat{f}(x_0) < \hat{f}(x_1)]$$

for x_0 sampled from $P(x \mid y = 0)$, and x_1 sampled from $P(x \mid y = 1)$

Precision-recall curve



Prediction error
vs
expected prediction error



Two scenarios

Two scenarios

A:

- Trained a model $\hat{f}_\tau(x)$ on a particular dataset τ
- Want to know, how well this particular $\hat{f}_\tau(x)$ will perform on new data

Two scenarios

A:

- Trained a model $\hat{f}_\tau(x)$ on a particular dataset τ
- Want to know, how well this particular $\hat{f}_\tau(x)$ will perform on new data

B:

- Chose a particular algorithm $\mathcal{A}: \tau \rightarrow \hat{f}_\tau$, for a particular problem – defined by the unknown population distribution $P(x, y)$
- Want to know, how well this algorithm performs on this problem

Two scenarios

A:

- Trained a model $\hat{f}_\tau(x)$ on a particular dataset τ
- Want to know, how well this particular $\hat{f}_\tau(x)$ will perform on new data
- **Prediction error:**

$$\text{Err}_\tau = \mathbb{E}_{x,y} \left[L \left(y, \hat{f}_\tau(x) \right) \right]$$

B:

- Chose a particular algorithm $\mathcal{A}: \tau \rightarrow \hat{f}_\tau$, for a particular problem – defined by the unknown population distribution $P(x, y)$
- Want to know, how well this algorithm performs on this problem

Two scenarios

A:

- Trained a model $\hat{f}_\tau(x)$ on a particular dataset τ
- Want to know, how well this particular $\hat{f}_\tau(x)$ will perform on new data
- **Prediction error:**

$$\text{Err}_\tau = \mathbb{E}_{x,y} \left[L \left(y, \hat{f}_\tau(x) \right) \right]$$

B:

- Chose a particular algorithm $\mathcal{A}: \tau \rightarrow \hat{f}_\tau$, for a particular problem – defined by the unknown population distribution $P(x, y)$
- Want to know, how well this algorithm performs on this problem
- **Expected prediction error:**

$$\text{Err} = \mathbb{E}_{x,y,\tau} \left[L \left(y, \hat{f}_\tau(x) \right) \right] = \mathbb{E}_\tau [\text{Err}_\tau]$$

Splitting to train and test



What kind of error do we estimate here?

Splitting to train and test



What kind of error do we estimate here?

How to estimate its variance?

Splitting to train, validation and test

When we do model selection, we use the left-out data to estimate the prediction error and minimize it (e.g., wrt the hyperparameters)

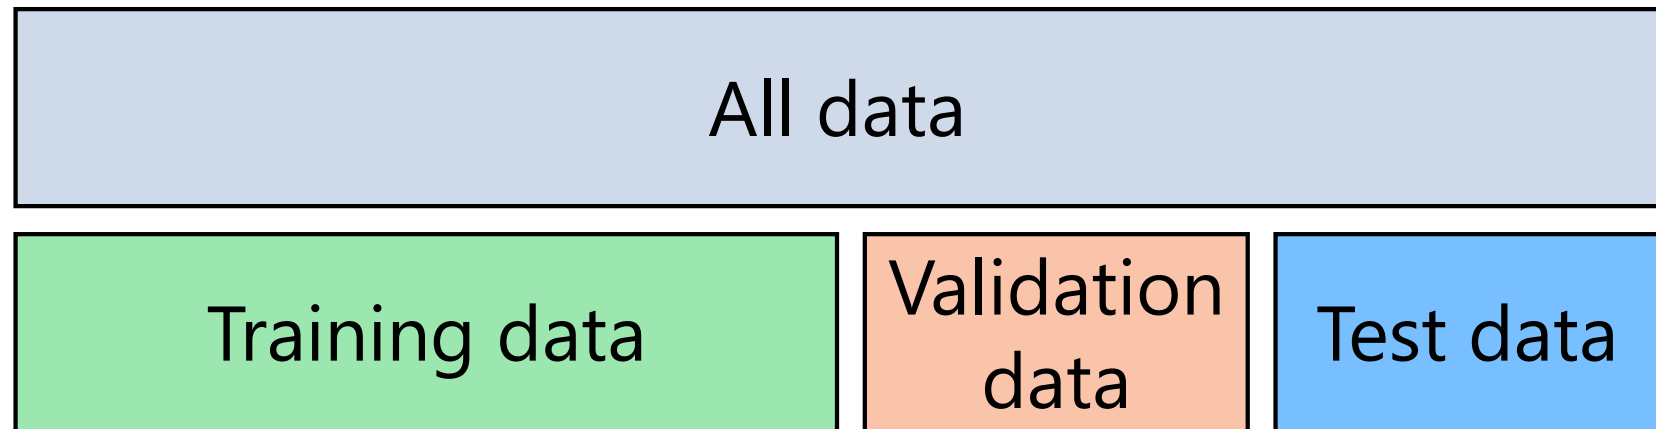
May 'overfit to test', so the resulting minimized error is not a good estimate of the prediction error

Splitting to train, validation and test

When we do model selection, we use the left-out data to estimate the prediction error and minimize it (e.g., wrt the hyperparameters)

May 'overfit to test', so the resulting minimized error is not a good estimate of the prediction error

Solution: tune on the **validation data**, do the final evaluation on the **test data**

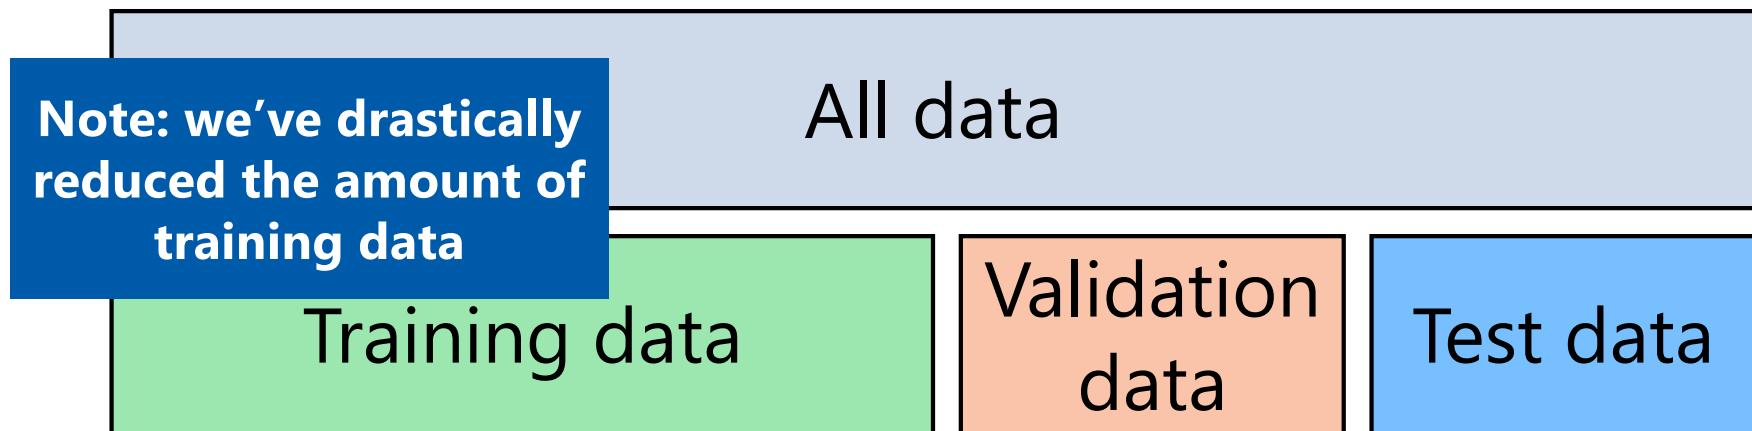


Splitting to train, validation and test

When we do model selection, we use the left-out data to estimate the prediction error and minimize it (e.g., wrt the hyperparameters)

May 'overfit to test', so the resulting minimized error is not a good estimate of the prediction error

Solution: tune on the **validation data**, do the final evaluation on the **test data**



Cross-validation



How to estimate the expected prediction error?

$$\text{Err} = \mathbb{E}_{x,y,\tau} \left[L \left(y, \hat{f}_{\tau}(x) \right) \right] = \mathbb{E}_{\tau} [\text{Err}_{\tau}]$$

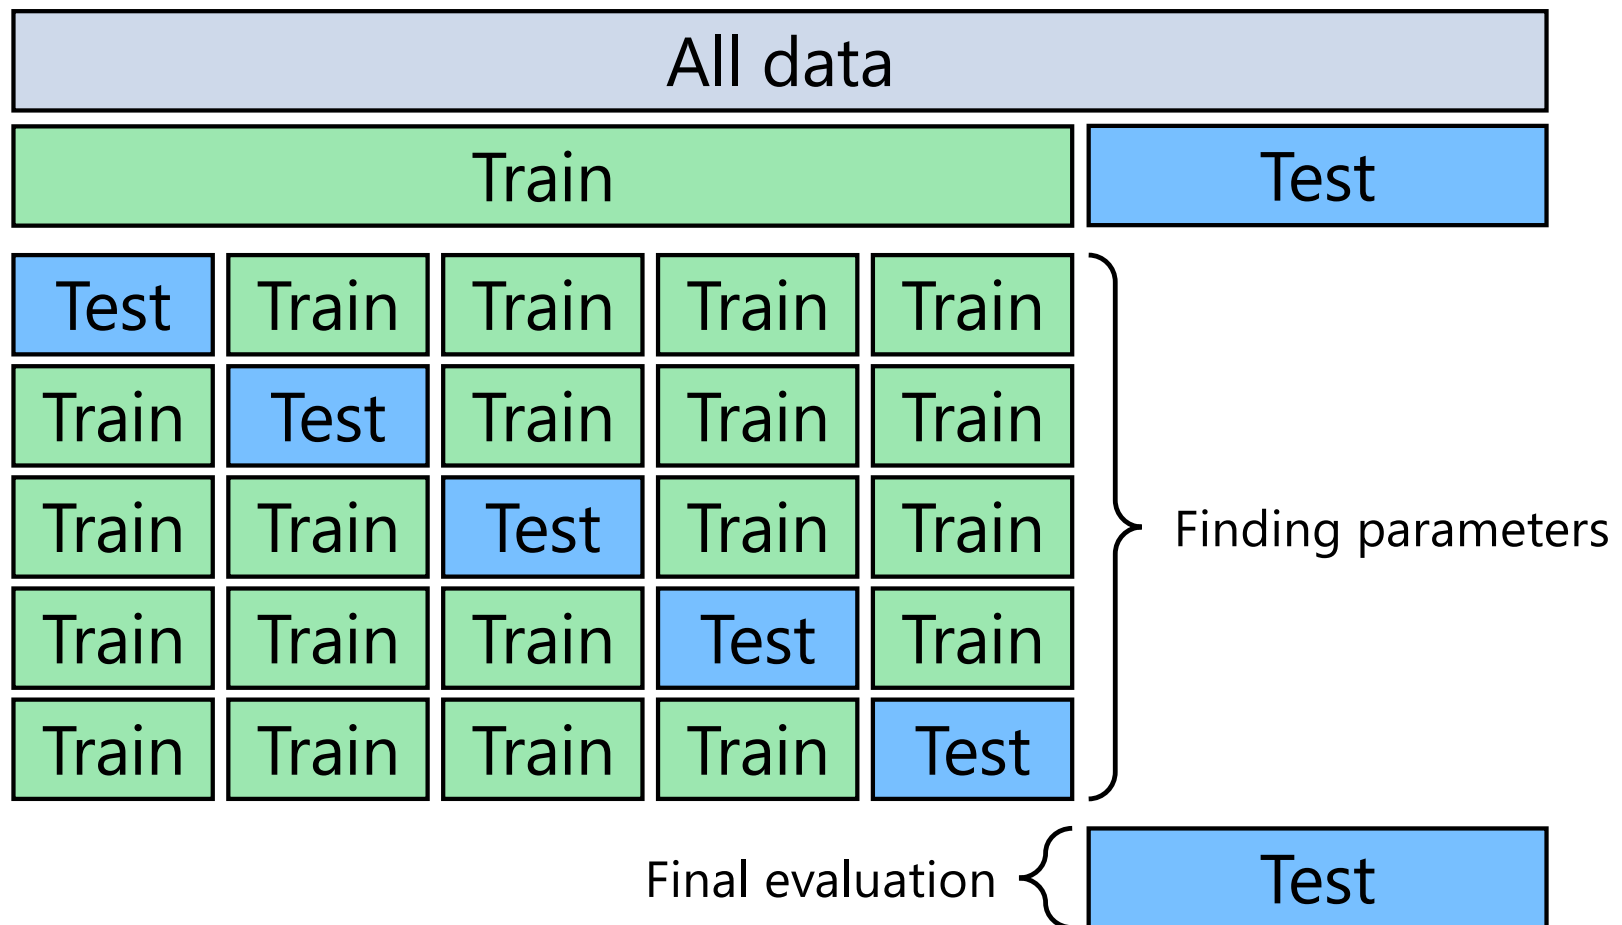
We can't just sample new training sets!
(unless there's **really** a lot of data available to us)

Note: Err_{τ} is by itself an estimate of Err ,
but since it's just a single observation we
know nothing about its variance

K-fold cross-validation

All data					
Test	Train	Train	Train	Train	Iteration 1
Train	Test	Train	Train	Train	Iteration 2
Train	Train	Test	Train	Train	...
Train	Train	Train	Test	Train	
Train	Train	Train	Train	Test	Iteration K

Hyperparameter tuning




Thank you!



Majid Sohrabi

msohrabi@hse.ru

Majid Sohrabi, NRU HSE



@MSOHRABI_CS