

Convolutional Networks

Motivation, working with images, trainable kernels

Machine Learning and Data Mining, 2024

Majid Sohrabi

National Research University Higher School of Economics



November 20, 2024

How to work with image-like data?



Working with images

Extremely high-dimensional input

- E.g. even a small 640x480 color image would make up almost 1M input features (pixel brightness levels in R, G and B)
- So a fully-connected hidden representation with just 100 units would require 100M parameters


Working with images

Extremely high-dimensional input

- E.g. even a small 640x480 color image would make up almost 1M input features (pixel brightness levels in R, G and B)
- So a fully-connected hidden representation with just 100 units would require 100M parameters

Is quite data-hungry to train when using fully-connected layers:


- Identifying an object on a picture would require examples with all possible locations of that object on the picture



Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object


Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object


Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object


Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object

Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object


Translational symmetry

A cat moved from one part of an image to another is still a cat


Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object


Translational symmetry

A cat moved from one part of an image to another is still a cat

Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest

0.001	0.003	0.876
0.000	0.002	0.001

Probability of locating the object

Translational symmetry

A cat moved from one part of an image to another is still a cat

Why don't we use the same model to look at different patches of an image trying to identify the object of interest:

This may be implemented with a 2D convolution!





Image
Input



Object of interest
**Convolutional
kernel**

0.001 0.003 0.876

0.000 0.002 0.001

Probability of locating the object
Output

2D convolution



2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

<i>a</i>	<i>b</i>	<i>c</i>
<i>d</i>	<i>e</i>	<i>f</i>
<i>g</i>	<i>h</i>	<i>i</i>

Input

<i>x</i>	<i>y</i>
<i>z</i>	<i>w</i>

Kernel

$ax + by + dz + ew$	

Output

2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

a	b	c
d	e	f
g	h	i

Input

x	y
z	w

Kernel

$ax + by$ $+ dz + ew$	$bx + cy$ $+ ez + fw$

Output

2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

a	b	c
d	e	f
g	h	i

Input

x	y
z	w

Kernel

$ax + by$ $+ dz + ew$	$bx + cy$ $+ ez + fw$
$dx + ey$ $+ gz + hw$	

Output

2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

a	b	c
d	e	f
g	h	i

Input

x	y
z	w

Kernel

$ax + by$ $+ dz + ew$	$bx + cy$ $+ ez + fw$
$dx + ey$ $+ gz + hw$	$ex + fy$ $+ hz + iw$

Output

2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

a	b	c
d	e	f
g	h	i

Input

x	y
z	w

Kernel

$ax + by$ $+ dz + ew$	$bx + cy$ $+ ez + fw$
$dx + ey$ $+ gz + hw$	$ex + fy$ $+ hz + iw$

Output

Different kernels may extract different features

Examples



Input

Blur:

$$\text{kernel} = \begin{pmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{pmatrix} \cdot \frac{1}{273}$$



Ouput

Examples



Input

Sharpen:

$$\text{kernel} = \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$



Ouput

Examples



Input

Edge detection:

$$\text{kernel} = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$



Output

Examples



Input

Edge detection:

$$\text{kernel} = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$



Ouput

In the context of deep learning, the kernel parameters are **trainable**

I.e. the network **learns the kernel** to **extract useful features**


2D convolution as a matrix multiplication

Unwrap the 2D images
into 1D vectors

Re-write the convolution
as a regular matrix-vector
multiplication

I.e. fully-connected layers
comprise convolutions

- Yet they are much more
complex



$ax + by + dz + ew$
$bx + cy + ez + fw$
$dx + ey + gz + hw$
$ex + fy + hz + iw$

=

x	y	0	z	w	0	0	0	0
0	x	y	0	z	w	0	0	0
0	0	0	x	y	0	z	w	0
0	0	0	0	x	y	0	z	w

×

a
b
c
d
e
f
g
h
i

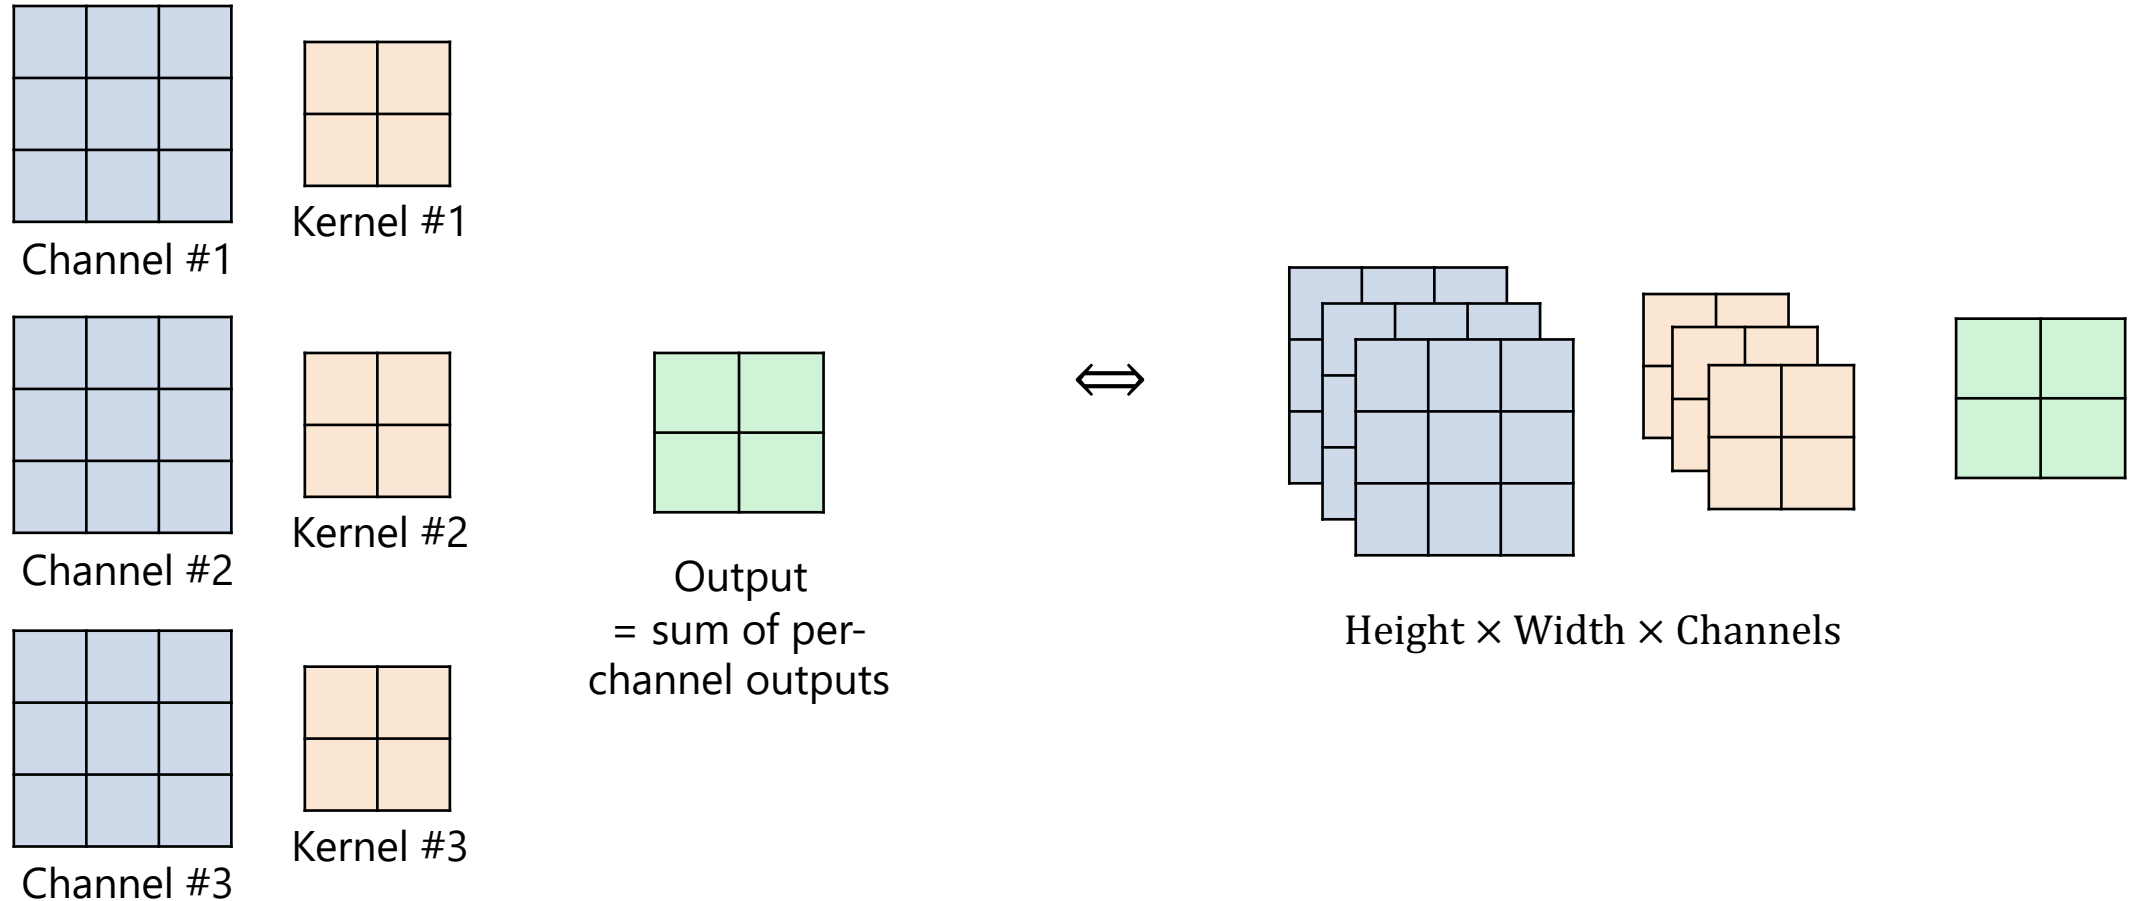
2D convolutional layers



Input channels

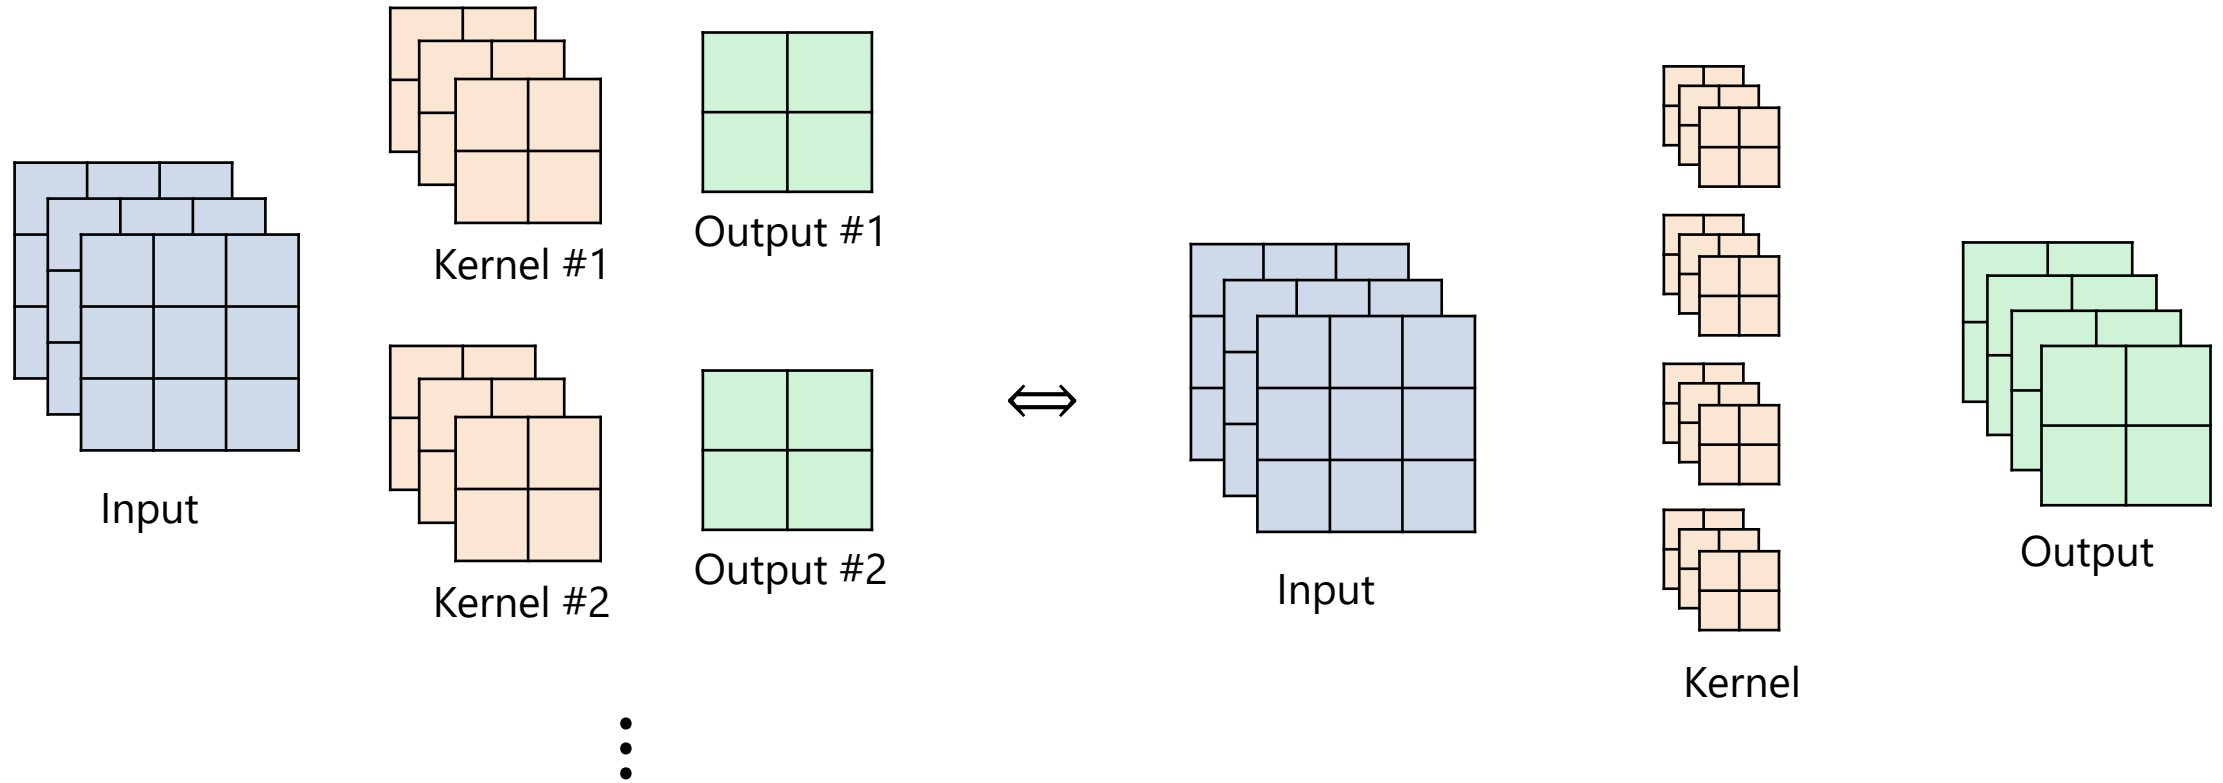
In practice images have multiple channels

- E.g. 3 color channels of a color image



Output channels

In practice we want to extract multiple features



Kernel becomes 4D: $H_K \times W_K \times C_{in} \times C_{out}$

Output becomes 3D: $H \times W \times C_{out}$

Putting it all together


$$\text{Output}(i, j, c_{out}) = \sum_{i', j', c_{in}} \text{Input}(i', j', c_{in}) \cdot \text{Kernel}(i' - i, j' - j, c_{in}, c_{out}) + \overset{\text{bias term}}{\downarrow} b_{out}$$

Note: with this approach the output width/height is smaller than the input width/height

- By how much (for a given kernel width and height)?

Sometimes the border of the input image is **padding** with some values (e.g. s.t. the output has the same size)

- Controlled by the “padding” parameter



Some other parameters:


- “stride” – by how many pixels the kernel window steps (equals 1 in the examples here)
- “dilation” – kernel “spread” (e.g. see [this animation](#))

Typical network architecture



Deep convolutional network

Image from: <https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f14>



Convolution is typically followed by a pooling operation (aggregating values of nearby pixels), e.g. maxpool:

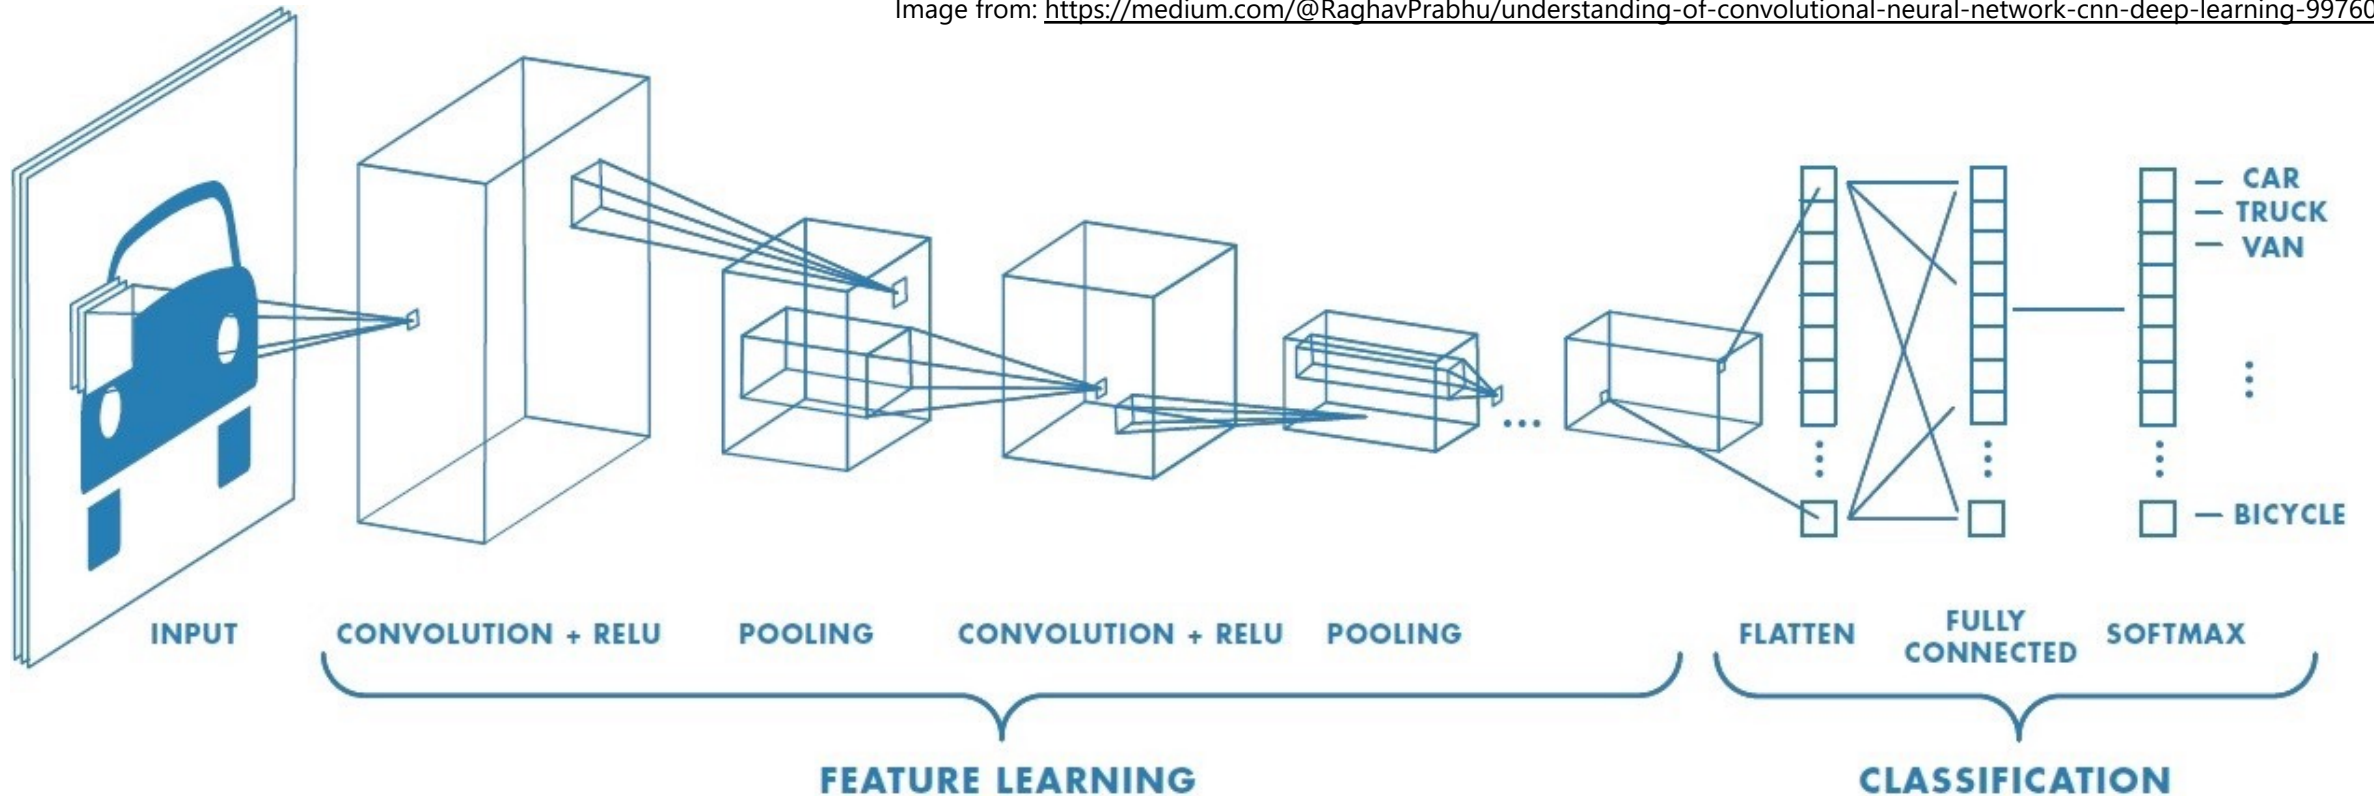
1	-2	5	7
2	1	4	-3
4	-9	2	0
-4	15	2	1

$\xrightarrow{\text{maxpool}(2 \times 2)}$

2	7
15	2

Deep convolutional network

Image from: <https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f14>



Convolution is typically followed by a pooling operation (aggregating values nearby pixels), e.g. maxpool:

Motivation:

Deeper layers extract "higher level" (i.e. more complex) features, and we're less interested in their spacial location.
Hence, fewer pixels and more channels.

1	5	2	0
-4	15	2	1


15	7
15	2

Thank you!

Majid Sohrabi



msohrabi@hse.ru



@MSOHRABI_CS